

Name:

Student USC ID:

Student USC email:

Row:

Seat:

Y

CSCI 561

Foundations of Artificial Intelligence

Exam #3 – Spring 2014

4:00-5:20pm

Friday, May 2, 2014

Instructor: Prof Tejada

(This exam is closed book, closed notes, closed everything.


No “cheat sheet” allowed.

No calculators, cell phones, or any electronic gadgets.)

Notes. As grading criteria, you will get full marks for each question that has the right answer (with appropriate reasoning or show of steps). You may be assigned a partial credit for question if deemed appropriate. If you face difficulty, or if things are unclear, write down your assumptions and proceed accordingly.

Reasoning about Uncertainty (30pts)

Bayesian Networks (20pts) Given this network calculate the following probabilities. Give both the formula and calculations with values. These questions are designed so that they can be answered with a minimum of computation. If you find yourself doing a copious amount of computation for each part, step back and consider whether there is simpler way to deduce the answer.



1. [4pts] $P(a, \neg b, c, \neg d)$

2. [4pts] $P(b)$

Bayesian Networks continued

3. [4pts] $P(a|b)$

4. [4pts] $P(d|a)$

5. [4pts] $P(d|a,c)$

Probability (10pts)

Circle **True** for each of the following statements if it is always true.

1. [2pts] [True/False] $P(A, B) = P(A)P(B)$

2. [2pts] [True/False] $P(A|B) = P(A)P(B)$

3. [2pts] [True/False] $P(A, B) = P(A)P(B) - P(A|B)$


4. [2pts] [True/False] $P(A, B, C) = P(A|B, C)P(B|C)P(C)$

5. [2pts] [True/False] If $P(A|B, C) = P(B|A, C)$ then $P(A|C) = P(B|C)$

Machine Learning (50pts)

Neural Networks (10pts)

In this network the input layer consists of units X_1 and X_2 . The hidden layer is Y_1 and Y_2 , and the output layer is Z_1 . Units Y_1 , Y_2 and Z_1 all have threshold functions where $t > 0$ for the unit to activate.



1. [1pt] When input $X_1 = 0$ and input $X_2 = 0$, what does the network output?
2. [1pt] When input $X_1 = 0$ and input $X_2 = 1$, what does the network output?
3. [1pt] When input $X_1 = 1$ and input $X_2 = 0$, what does the network output?
4. [1pt] When input $X_1 = 1$ and input $X_2 = 1$, what does the network output?
5. [3pts] Can a weight be changed for the network to compute the Boolean function OR? If yes, which weight and what value should it be changed to?
6. [3pts] Can a single weight be changed for the network to compute the Boolean function XOR? If yes, which weight and what value should it be changed to?

Candidate Elimination (20pts)

Reward Card Domain: This domain is a deck of 52 cards (13 cards of each suit ♠ ♦ ♥ ♣). A subset of the deck are considered “Reward” cards. The goal is given positive and negative examples of Reward cards learn the concept that represents the Reward cards only.

$(r=1) \vee \dots \vee (r=10) \vee (r=J) \vee (r=Q) \vee (r=K) \Leftrightarrow \text{ANY-RANK}(r)$

$(r=1) \vee \dots \vee (r=10) \Leftrightarrow \text{NUM}(r)$

$(r=J) \vee (r=Q) \vee (r=K) \Leftrightarrow \text{FACE}(r)$

$(s=\spadesuit) \vee (s=\clubsuit) \vee (s=\diamondsuit) \vee (s=\heartsuit) \Leftrightarrow \text{ANY-SUIT}(s)$


$(s=\spadesuit) \vee (s=\clubsuit) \Leftrightarrow \text{BLACK}(s)$

$(s=\diamondsuit) \vee (s=\heartsuit) \Leftrightarrow \text{RED}(s)$

A hypothesis is represented by rs, with:

- $r \in \{a, n, f, 1, \dots, 10, j, q, k\}$, where a=ANY-RANK, n=NUM, f=FACE
- $s \in \{a, b, r, \clubsuit, \spadesuit, \diamondsuit, \heartsuit\}$, where a=ANY-SUIT, b=BLACK, r=RED

The hypothesis space for the Reward Card domain can be seen as a partial order, where G is the set of most general hypotheses (aa = ANY-RANK, ANY-SUIT), and S is the set of most specific hypotheses (4♣). Given this subset of the partial order, answer the questions below.



1. [2pts] Given 4♣ is a positive example of a Reward card, list the hypotheses that would classify 4♣ as a true negative.

Candidate Elimination continued

-
2. [2pts] Given $7\clubsuit$ is a positive example of a Reward card, list the hypotheses that would classify $7\clubsuit$ as a true positive.
3. [2pts] Given $5\heartsuit$ is a negative example of a Reward card, list the hypotheses that would classify $5\heartsuit$ as a false positive.

Candidate Elimination continued

4. [2pts] Given $2\spadesuit$ is a positive example of a Reward card, list the hypotheses that would classify $2\spadesuit$ as a false negative.

5. [4pts] Given this set of training examples for Reward card, ($4\clubsuit$ is a positive example, $7\clubsuit$ is a positive example, $5\heartsuit$ is a negative example, $2\spadesuit$ is a positive example), list the hypotheses that would correctly classify all examples.

Candidate Elimination continued

6. [4pts] The example $j\spadesuit$ is added to the training set for Reward card. If it is a positive example, list the hypotheses that would correctly classify all examples. How many of the 52 cards are Reward cards?
7. [4pts] The example $j\spadesuit$ is added to the training set for Reward card. If it is a negative example, list the hypotheses that would correctly classify all examples. How many of the 52 cards are Reward cards?

Reinforcement Learning (20pts)

Given a Gridworld domain, where terminal states (1,3), (4,3), and (4,2) have rewards 50, 500, and -50 respectively, the set of possible actions are {N,E,S,W, or X for terminal states}, the agent moves deterministically, all V and Q values for non terminal states have been initialized to 0, answer the questions below.

3	50			500
2				-50
1				
	1	2	3	4

Circle the letter that corresponds to the best answer for the question:

1. [4pts] What are the optimal values, V^* of each state in the above grid if $\gamma = 0.5$, $c(a)=0$, $R(s)=0$ for non terminal states?

(Remember $V_{t+1}(s) = R(s) + \text{Max}_{a \in A} \{c(a) + \gamma \sum_{s' \in S} \text{Pr}(s'|a,s) V_t(s')\}$)

- a. $V_{(1,1)}=15.75$, $V_{(1,2)}=25$, $V_{(2,1)}=31.25$, $V_{(2,3)}=125$, $V_{(3,1)}=62.5$, $V_{(3,2)}=125$, $V_{(3,3)}=250$, $V_{(4,1)}=25$
- b. $V_{(1,1)}=12.5$, $V_{(1,2)}=25$, $V_{(2,1)}=31.25$, $V_{(2,3)}=125$, $V_{(3,1)}=62.5$, $V_{(3,2)}=125$, $V_{(3,3)}=250$, $V_{(4,1)}=31.25$
- c. $V_{(1,1)}=15.75$, $V_{(1,2)}=25$, $V_{(2,1)}=31.25$, $V_{(2,3)}=125$, $V_{(3,1)}=62.5$, $V_{(3,2)}=125$, $V_{(3,3)}=250$, $V_{(4,1)}=31.25$
- d. $V_{(1,1)}=12.5$, $V_{(1,2)}=25$, $V_{(2,1)}=25$, $V_{(2,3)}=25$, $V_{(3,1)}=50$, $V_{(3,2)}=100$, $V_{(3,3)}=250$, $V_{(4,1)}=25$
- e. None of the above

Reinforcement Learning (continued)

2. [4pts] What are the Q values of state (3,2) in the above grid if $\gamma = 0.5$, $c(a)=0$, $R(s)=-2$ for non terminal states?

(Remember $Q_{t+1}(a,s) = R(s) + c(a) + \gamma \sum_{s'} \Pr(s'|a,s) Q_t(a',s')$)

- a. $Q_{((3,2),N)}=120$, $Q_{((3,2),E)}=-27$, $Q_{((3,2),S)}=58$
- b. $Q_{((3,2),N)}=120$, $Q_{((3,2),E)}=-27$, $Q_{((3,2),S)}=27$
- c. $Q_{((3,2),N)}=125$, $Q_{((3,2),E)}=-25$, $Q_{((3,2),S)}=62.5$
- d. $Q_{((3,2),N)}=120$, $Q_{((3,2),E)}=-27$, $Q_{((3,2),S)}=31.5$
- e. None of the above

For questions 3-5:

When the agent is not given the map with the reward function, it starts Q-learning with these 3 training examples or episodes/epochs:

E1: (1,1), N, 0, (1,2), N, 0, (1,3), X, 50

E2: (1,1), N, 0, (1,2), N, 0, (1,3), X, 50

E3: (1,1), N, 0, (1,2), S, 0, (1,1), E, 0, (2,1), E, 0, (3,1), E, 0, (4,1), N, 0, (4,2), X, -50

During Q-Learning each example is processed in the order the states are observed, for example in E1 the Q values are calculated for the observed states in this order (1,1), (1,2), (1,3).

Remember that the Q-Learning equation is:

$$Q(a,s) \leftarrow Q(a,s) + \alpha(R(s) + \gamma \text{Max}_{a' \in A} Q(a',s') - Q(a,s))$$

3. [4pts] Given $\alpha=1$ and $\gamma = 0.5$, what are the non-zero Q values after observing all 3 of these examples?

- a. $Q_{((1,1),N)}=50$, $Q_{((1,2),N)}=50$, $Q_{((1,3),X)}=50$, $Q_{((4,2),X)}=-50$
- b. $Q_{((1,1),N)}=12.5$, $Q_{((1,2),N)}=25$, $Q_{((1,3),X)}=50$, $Q_{((4,2),X)}=-50$
- c. $Q_{((1,1),N)}=12.5$, $Q_{((1,2),S)}=6.25$, $Q_{((1,2),N)}=25$, $Q_{((1,3),X)}=50$, $Q_{((4,2),X)}=-50$
- d. $Q_{((1,1),N)}=12.5$, $Q_{((1,2),S)}=12.5$, $Q_{((1,2),N)}=25$, $Q_{((1,3),X)}=50$, $Q_{((4,2),X)}=-50$
- e. None of the above

Reinforcement Learning (continued)

3	50	E=-0 W=0	W=0 E=250 S=125	500
2	N=25 S= 7.875		N=125 E=-25 S=31.5	-50
1	N=12.5 E=15.75	E=31.25 W=7.875	N=62.5 W=3.9375 E= -12.5	N=-25 W=31.5
	1	2	3	4

4. [4pts] After Q-Learning with many more examples, the Q values are shown in the grid above. What is the optimal policy given these values?
- (1,1) N, (1,2) N, (2,1) E, (3,1) N, (4,1) W, (3,2) N, (3,3) E
 - (1,1) E, (1,2) S, (2,1) E, (3,1) N, (4,1) W, (3,2) N, (3,3) E
 - (1,1) E, (1,2) N, (2,1) E, (3,1) N, (4,1) W, (3,2) N, (3,3) E
 - All of the above
 - None of the above
5. [4pts] If now there is a “living reward” of -5 for all non terminal states ($R(s) = -5$), so that all new examples would see a -5 instead of 0 at non terminal states. How would this affect the agents policy?
- It would prefer actions providing the shortest path to state (4,2)
 - It would prefer actions providing the shortest path to state (4,3)
 - It would prefer actions providing the shortest path to state (1,2)
 - All of the above
 - None of the above

Short Answer (20pts)

Keep your answers brief, one or two sentences.

1. [4pts] List 2 reasons why machine learning is needed?
2. [4pts] What is Ockham's razor? How is it used in decision tree learning?
3. [4pts] What does the size of a hypothesis class have to do with generalization and overfitting?

4. [4pts] In reinforcement learning, why can it be useful to sometimes act in a way that is believed to be suboptimal?

5. [4pts] For the questions below choose 2 research topics from Lecture 12 on Communicating, Perceiving, and Acting and describe each topic, including the limitations/pitfalls and future directions.