

Total score **34**

34A2C7D4-EA89-4F9B-AF3D-DCBE02377962

csci561-su17-exam3

#22 2 of 18



Q1TF

6

1. [20%] Machine Learning & General AI Knowledge

True or False [10%]: For each of the statements below, fill in the bubble T if the statement is always and unconditionally true, or fill in the bubble F if it is always false, sometimes false, or just does not make sense.

1)	<input checked="" type="checkbox"/>	F
2)	<input checked="" type="checkbox"/>	F
3)	<input checked="" type="checkbox"/>	F
4)	<input checked="" type="checkbox"/>	F
5)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
6)	T	<input checked="" type="checkbox"/>
7)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
8)	T	<input checked="" type="checkbox"/>
9)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
10)	<input checked="" type="checkbox"/>	F

1) [1%] The ROC curve shows the tradeoff between sensitivity and specificity.

2) [1%] Specificity quantifies the avoiding of false negatives.

3) [1%] Random Forest is an ensemble method combining decision trees.

4) [1%] Cross validation helps to avoid overfitting.

5) [1%] Unsupervised learning is learning from unlabeled data to maximize a cumulative reward.

this is reinforcement learning (unsupervised
l. has no reward signal)

6) [1%] The more layers a neural network has the less likely it is to overfit.

7) [1%] Support vectors are data points that lie closest to the hyperplane.

8) [1%] A highly accurate classifier has an ROC curve that is very close to the diagonal line.

9) [1%] K-nearest neighbors is a discriminative classification algorithm.

KNN is a generative model.

10) [1%] A decision tree can represent any Boolean function.

90B60460-BF42-4791-87F9-987C03EC72C4

csci561-su17-exam3

#22 3 of 18




Q1MC

1

Multiple Choice [10%]

11) [5%] Consider the following Constraint Satisfaction Problem (CSP) graph.

Each of the variables A, B, C, D, E, F can take values 1, 2, 3, 4. The constraints are inequality relationships between two variables, as depicted in the graph.



Which of the following statements are true?

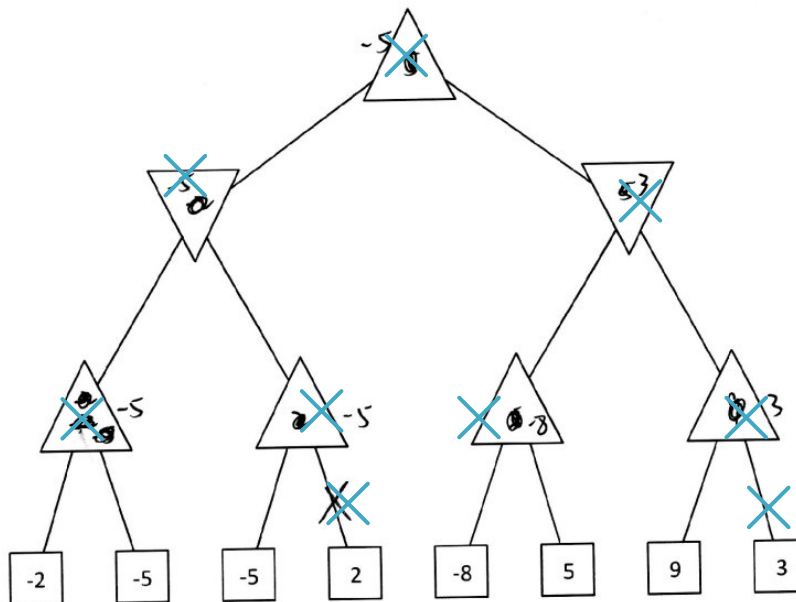
- 1) ☒ D is the only most constrained variable.
- 2) ☒ D is the only most constraining variable.
- 3) ☒ If only A and B are assigned, then D has the least number of remaining values.
- 4) ☒ If A=1, B=2, C=3 then the only remaining values for E are 1, 2, 3 after forward-checking.
- 5) ☒ If the domain size was 3 the CSP could still be solved.

1) There is no assignment so far, so all variables have the same remaining values 1, 2, 3, 4.

4) This would be the case for arc-consistency checking.

12) [5%] Complete the following Minimax tree by filling in the min/max values to the empty nodes.

Indicate which edges get pruned by Alpha-Beta pruning by striking through them:



3 / 18

E5FD2F56-6164-443E-BE5D-83DA7AD5349C

csci561-su17-exam3

#22 4 of 18



Q2(No) Not graded

2. [20%] Inference in First Order Logic

Consider the following 13 facts and rules that are added to a knowledge base: *Connected*, *Speed*, *Distance*, *Sum*, *Quotient* and *Time* are predicates, *USC*, *LAX*, *SFO*, *GGB*, *Car*, *Metro* and *Plane* are constants, and $x, y, s, g, t, pt, rt, pg, m, v$ and d are variables that are universally quantified. *Connected*(x, y, z) means place x and place y could be connected by transportation tool z . *Speed*(x, y) means the speed of transportation x is y . *Distance*(x, y, z) means the distance between place x and place y is z . *Sum*(x, y, z) means sum of x and y is z . *Quotient*(x, y, z) means quotient of x divided by y is z . *Time*(x, y, z) means the time needed to travel from place x to place y is z .

1. *Connected*(*USC*, *LAX*, *Car*)
2. *Connected*(*LAX*, *SFO*, *Plane*)
3. *Connected*(*SFO*, *GGB*, *Metro*)
4. *Speed*(*Car*, 50)
5. *Speed*(*Metro*, 40)
6. *Speed*(*Plane*, 400)
7. *Distance*(*USC*, *LAX*, 25)
8. *Distance*(*LAX*, *SFO*, 400)
9. *Distance*(*SFO*, *GGB*, 20)
10. *Sum*(x, y, z)
11. *Quotient*(x, y, z)
12. *Time*($x, x, 0$)
13. $Connected(s, pg, m) \wedge Speed(m, v) \wedge Distance(s, pg, d)$
 $\wedge Quotient(d, v, pt) \wedge Time(pg, g, rt) \wedge Sum(pt, rt, t)$
 $\Rightarrow Time(s, g, t)$

7C4ECC6A-B491-41DB-921E-F800E553C3BC

csci561-su17-exam3

#22 5 of 18




Q2

0

[20%] Show how backward chaining can be used to infer whether it takes 2 hours to reach the Golden Gate Bridge (GGB) from USC (i.e., $\text{Time}(\text{USC}, \text{GGB}, 2)$). Report each step of the inference in detail in terms of unifications, rule firings, etc.

1. 4, 7, 10, 11, 12, 13 $\Rightarrow \text{Time}(\text{USC}, \text{LAX}, 0.5)$ (14)
2. 6, 8, 10, 11, 13, 14 $\Rightarrow \text{Time}(\text{USC}, \text{SFO}, 1.5)$ (15)
3. 5, 9, 10, 11, 13, 15 $\Rightarrow \text{Time}(\text{USC}, \text{GGB}, 2)$ (16)



5 / 18

9CCE3D70-1859-4648-BD8E-D6A6B752CE21

csci561-su17-exam3

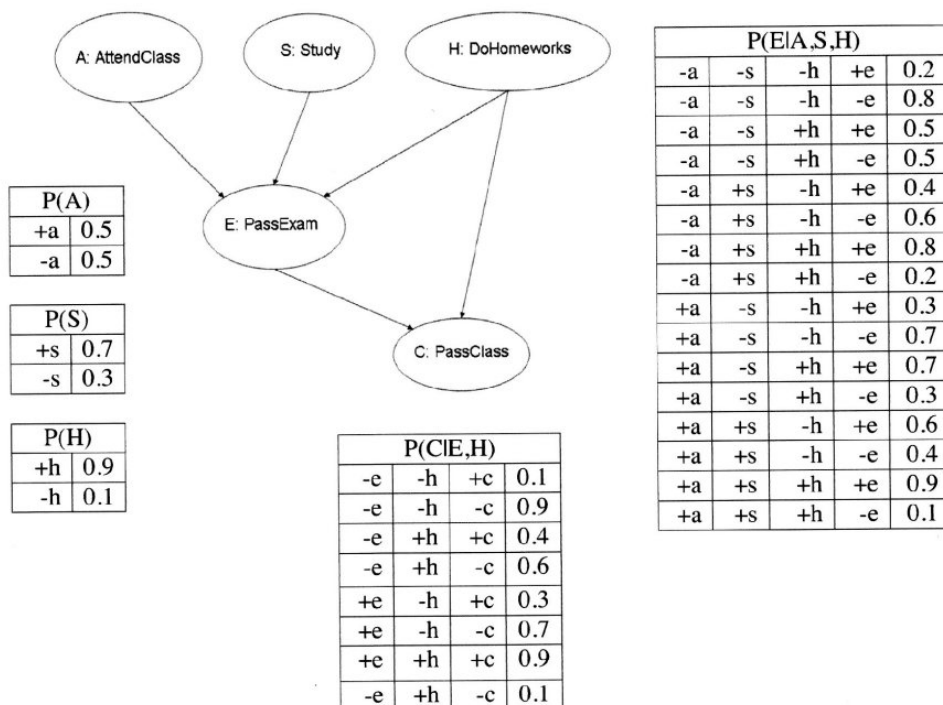
#22 6 of 18



Q3 Not graded

3. [15%] Bayesian Network.

Suppose that a student PassExam(E), could be caused by AttendClass(A), Study(S), DoHomeworks(H). PassClass(C) could be caused by PassExam(E) and DoHomeworks(H). As long as you have the formula correct, you will get full credit. There is no need to calculate the final numbers.



6 / 18

0F91744C-9D54-4F87-96BB-913CE311118A

csci561-sul7-exam3

#22 7 of 18



Q3AB 6

3A. [3%] Compute the following entry from joint distribution: $P(+a, +s, +h, +e, +c)$.

$$P(+a, +s, +h, +e, +c) = P(+a) \cdot P(+s) \cdot P(+h) \cdot P(+e|+a, +s, +h) \cdot P(+c|+e, +h)$$

3B. [3%] What is the probability of passing the class, given that you attended class and studied, but didn't do the homework?

$$\begin{aligned}
 P(+c|+a, +s, -h) &= \frac{P(+c, +a, +s, -h)}{P(+a, +s, -h)} \\
 &= \frac{\sum_{e \in \{+e, -e\}} P(+a) \cdot P(+s) \cdot P(-h) \cdot P(e|+a, +s, -h) \cdot P(+c|e, -h)}{P(+a) \cdot P(+s) \cdot P(-h)} \\
 &= \frac{P(+e|+a, +s, -h) \cdot P(+c|+e, -h) + P(-e|+a, +s, -h) \cdot P(+c|-e, -h)}{1}
 \end{aligned}$$

7 / 18

6CC2D872-51CF-4D46-95D0-DFEE07B2D1F6

csci561-su17-exam3

#22 8 of 18



Q3C 2

3C. [3%] Compute $P(+a \mid +c, +h)$.

$$\begin{aligned}
 P(+a \mid +c, +h) &= \frac{P(+a, +c, +h)}{P(+c, +h)} \\
 &= \frac{\sum_{E, S} P(+a) \cdot P(+h) \cdot P(E \mid +a, S, +h) \cdot P(S) \cdot P(+c \mid E, +h)}{\sum_{A, S, E} P(+h) \cdot P(+c \mid E, +h) \cdot P(E \mid A, S, +h) \cdot P(A) \cdot P(S)}
 \end{aligned}$$

8 / 18

3912F61E-4541-437B-9FE2-EDB6DA612C77

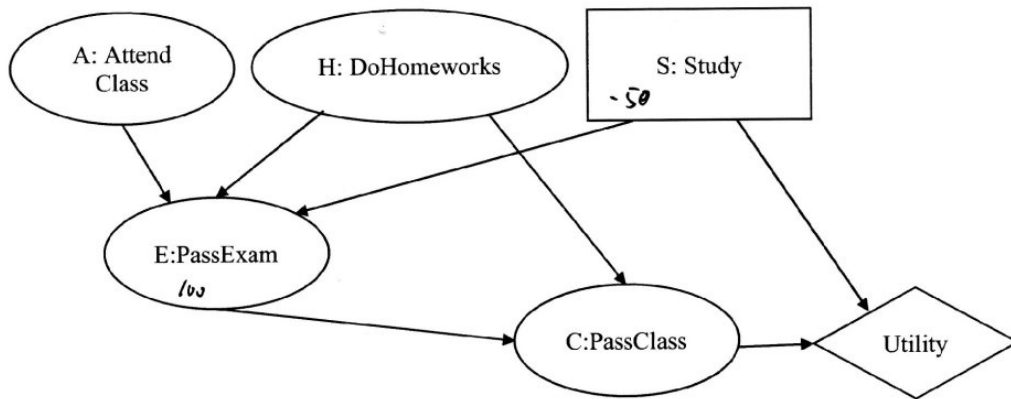
csci561-su17-exam3

#22 9 of 18



Q3D 0

3D. [6%] Now, consider a student who has the choice to study (Study(S)) or not study for passing the exam. We will model this as a decision problem with one Boolean decision node, S, indicating whether the agent chooses to study. The remaining nodes are chance nodes. The probabilities are the same as above. There is also a utility node U. We have the following utility function: the utility for PassClass is 100 and the utility for Not PassClass is -10. The utility for Study is -50 and the utility for Not Study is 0. Calculate the MEU for decision node S.



9 / 18

DD9366B5-8F0B-4373-827E-A2975ABBFBE2

csci561-su17-exam3

#22 10 of 18

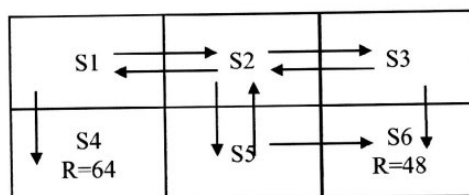


Q4A

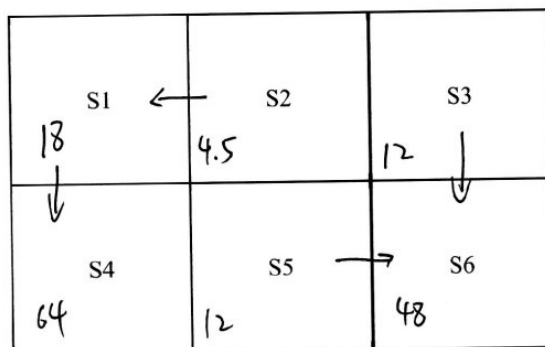
3

4. [20%] Markov Decision Process and Reinforcement Learning

Consider the 6-state Markov Decision process below. The goals with rewards are in state S4 and S6. At each state, the possible transitions are deterministic and indicated by the arrows. You get a reward of $R_4=64$ if you get to the goal S4 and a reward of $R_6=48$ if you get to the goal S6. In this problem, you have to get the final answer correct to get credits.



4A. [4%] Consider a discount factor of $\gamma = \frac{1}{4}$. On the figure below, show the optimal value V^* for each state and the arrows corresponding to the set of the optimal actions.



F4951BA9-278B-4CCD-97A5-57AA09C63A7F

csci561-su17-exam3

#22 11 of 18



Q4BC

9

4B. [5%] What values of γ would result in a different optimal action in S3? Indicate which policy action changes.

~~$48\gamma^2 > 48\gamma$~~

$$64\gamma^3 > 48\gamma$$

$$\gamma^2 > \frac{48}{64} = \frac{3}{4}$$

$$\gamma > \frac{\sqrt{3}}{2}$$

If $\gamma > \frac{\sqrt{3}}{2}$, $Q(S_3, \text{left}) > Q(S_3, \text{down})$, so S_3 will choose a different optimal, going left.

4C. [8%] Suppose the action is no longer deterministic. The discount factor is still $\gamma = \frac{1}{4}$. Each action has a failure probability $f=0.2$, i.e. the action is now stochastic. Once an action is taken, with probability 0.2, it would stay in the original grid. On the figure below, show the optimal value V^* for each state and the arrows corresponding to the set of the optimal actions.

$$Q = (0.8 \times 64 + 0.2 \times \gamma) \gamma$$

$$4X = 51.2 + 0.2X$$

$$3.8X = 51.2$$

$$X = \frac{51.2}{3.8}$$

$$38 \overline{) 512}$$

$$\underline{38}$$

$$132$$

$$\underline{114}$$

$$180$$

$$\underline{152}$$

$$28$$

$$4X = 0.8 \times 13.5 + 0.2X$$

$$3.8X = 10.8$$

$$X = \frac{10.8}{3.8}$$

$$38 \overline{) 108}$$

$$\underline{76}$$

$$320$$

S1	S2	S3
13.5 ↓	2.8	10.1 ↓
S4	S5	S6
64	10.1	48

$$X = \frac{1}{4} (0.8 \times 48 + 0.2X)$$

$$4X = 38.4 + 0.2X$$

$$3.8X = 38.4$$

$$X = \frac{38.4}{3.8}$$

$$38 \overline{) 384}$$

$$\underline{38}$$

$$20$$

$$24.8$$

11 / 18

EFC31AA5-FBDA-4CA6-9E25-6C72BDED205A

csci561-su17-exam3

#22 12 of 18



Q4D 3

4D. [3%] In this scenario, the discount factor is $\gamma = 1$ and the learning rate $\alpha = 0.5$. S6 has the reward 10 and S4 has the reward -10. All other states have reward 0. Now we no longer know the details of the transition probabilities ahead of time. We must instead use reinforcement learning to compute the necessary values. We have the following sequence of actions.

Remember that $Q(s, a)$ is initialized with 0 for all (s, a) .

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a))$$

Example 1			
s	a	s'	r
S1	Right	S2	0
S2	Down	S5	0
S5	Right	S6	0
S6	Exit	Terminal	10

Example 2			
s	a	s'	r
S1	Right	S2	0
S2	Down	S5	0
S5	Right	S6	0
S6	Exit	Terminal	10

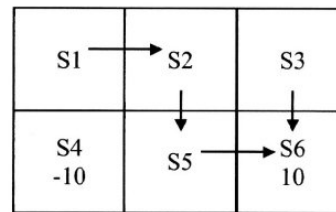
$$5 + \frac{1}{2} \cdot 10 = 7.5$$

Q = 5

$$0.5 \times 10$$

5

Example 3			
s	a	s'	r
S1	Right	S2	0
S2	Down	S5	0
S5	Right	S4	0
S4	Exit	Terminal	-10



5-

After 3 example runs, what is the Q-value of (S5, Right)?

After E2:

$$Q(S5, \text{Right}) = 0.5 \times 10 = 5$$

So, After E3:

$$\Rightarrow Q(S5, \text{Right}) = 7.5$$

12 / 18

DC5FB8C8-0E3F-4FA8-80BF-946B960EAAD8

csci561-su17-exam3

#22 13 of 18



Q5 Not graded

5. [15%] Decision Trees

Assume we want to train a machine to decide whether someone takes an Uber or just walks based on weather information and the length of the path. Our training data is provided below. The weather outlook can be either sunny (s), cloudy (c) or rainy (r). The humidity level can be either high (h) or normal (n). The destination is either near (n) or far (f).

Trip	Outlook	Humidity	Destination	TakeUber?
T1	s	h	n	y
T2	s	h	f	y
T3	c	h	n	n
T4	r	h	n	n
T5	r	n	n	n
T6	r	n	f	n
T7	c	n	f	n
T8	s	h	n	y
T9	s	n	n	n
T10	r	n	n	n
T11	s	n	f	n
T12	c	h	f	n
T13	c	n	n	n
T14	r	h	f	y
T15	r	n	n	n
T16	c	h	f	n

55,

CFA4CACC-C828-4500-9DBB-66072C8A5D4B

csci561-su17-exam3

#22 14 of 18



Q5A 3

5A. [6%] Which attribute would be at the root of the tree? Give the reason for your answer in terms of entropy and information gain (no need for exact calculation of the expressions).

~~Outlook: 1f~~

Outlook: 5s, 5c, 8r, $1(\frac{5}{16}, \frac{5}{16}, \frac{3}{16})$

Humidity: 8h, 8n, $1(\frac{3}{16}, \frac{3}{16})$

Destination: 7f, 9n, $2(\frac{7}{16}, \frac{9}{16})$

Take Uba?: 4y, 12n, $1(\frac{4}{16}, \frac{12}{16})$

For humidity: entropy $I_H = 1$, larger than others.

I_G is positive, when humidity is the root.

⇒ Humidity would be at the root.

Missing correct entropy expressions [-3%]

14 / 18

90D25E11-20C8-4064-9410-6E7B29F60AE4


csci561-su17-exam3

#22 15 of 18



Q5B 0

5B. [9%] Draw the learned decision tree and write down the learned Boolean expression for TakeUber as a disjunction of clauses.



Please refer to the solution. [-9%]

15 / 18

29D76AC4-87D9-44F5-B95C-DD73E84442F7

csci561-su17-exam3

#22 16 of 18



Q6A

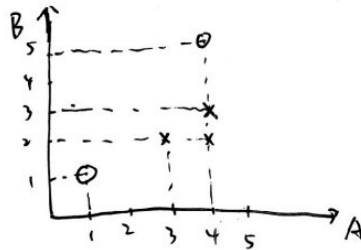
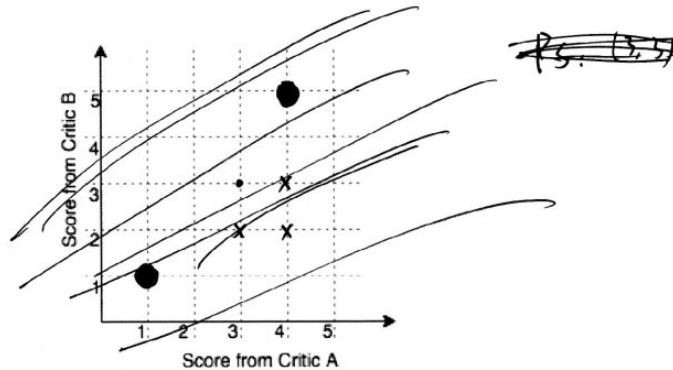
1

6. [10%] Neural Nets

Assume you want to predict how well a show will do with the general audience, based on the scores of two critics who score the show on the scale of 1 to 5. Here are five data points from some previous shows, including the critics' score and the performance of the shows:

Show	Critic A Score	Critic B Score	Did the audience like the show?
1	1	1	No
2	3	2	Yes
3	4	5	No
4	4	3	Yes
5	4	2	Yes

6A. [2%] Determine if the data is linearly separable by plotting it on the 2D plane below. Represent positive data points with x's and negative data points with o's.



16 / 18

52715B68-2917-479D-A16E-349E2E80838F

csci561-su17-exam3

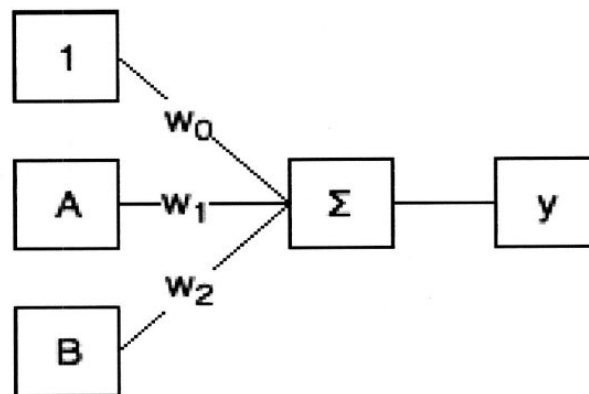
#22 17 of 18



Q6B(No) Not graded

6B. [8%] After establishing linear separability, you decide to use a perceptron to classify the data, using the scores as features. In the figure below, A, B are inputs from each critic and constant input 1 is a bias input, y is output.

So, the perceptron would be like this (If $w_0 + w_1*A + w_2*B > 0$, the audience will like the show the output ($y = 1$); otherwise they won't ($y = -1$):



This is how perceptron update works:

Start with an initial vector of weights.

For each training instance (y is the output and y^* is the expected result):

If correct ($y = y^*$), no update is needed.

If wrong: update the weight vector by adding or subtracting the input vector.

Subtract if y^* is -1.

For example, for the input vector $\langle 1, A, B \rangle$, if at a training iteration the input is (1, 2, 1), the weight vector is (1, 2, 3), the output is 1, and the expected result is -1, the updated weight vector will be: (0, 0, 2).

BA041A0E-D83D-4052-8B5C-7C9E6D5F0004

csci561-su17-exam3

#22 18 of 18



Q6B

0

Assume we start with the weights $\{w_0: 0, w_1: 0, w_2: 0\}$ (weight vector $(0, 0, 0)$). Determine the weights after two updates. Calculate the accuracy of the perceptron on the data after these two updates.

Input $(1, 2, 1)$

~~Let $w = (-1, 2, 1)$~~

$w = (0, 0, 0)$

output = -1

1st ~~$w = (-1, 2, 1)$~~ $w = (0, 0, 0)$

~~output = -1~~

2nd ~~$w = (-1, 2, 1)$~~ $w = (0, 0, 0)$

Accuracy = 1

18 / 18